

Global Seller Center UX Blueprint — v2

PRODUCT STRATEGY • UX ARCHITECTURE • VALIDATION PLAN

An evidence-aware, low-training operating system for marketplace sellers across regions, devices, currencies, and fulfillment models.

Global core • Market configuration • Local adapters

Revision 2 — August 2026. Revised per internal review: adds an explicit V1 bet/cut list, splits the architecture into a minimum viable core versus deferred capabilities with a documented-debt register, demotes contribution margin to the optimization phase, and hardens Stage 1 discovery into a formal go/no-go gate.

Section 01 — Executive decision

Keep the operational ideas. Correct the scientific framing. Build one global product with market-specific rules — **and ship a deliberately narrow v1.**

DECISION: Proceed with the seller-center concept as a product hypothesis, not as validated research. The strongest ideas are progressive listing, print-first fulfillment, explicit financial breakdowns, contextual messaging, inline inventory editing, and configurable payouts. The original psychology labels are not treated as proof of effectiveness.

The revised direction

Keep	Correct	Add	Cut from v1
Action-first dashboard; rapid listing; batch fulfillment; financial ledger; unified order context; inline stock editing.	Replace "net profit" with estimated proceeds; remove "guaranteed cash"; stop using Peak-End Rule to justify payouts.	Market-configuration layer, accessibility, low-bandwidth recovery, permission controls, versioned fee rules, a real validation program.	Contribution margin view, predictive prioritization, multi-market configuration console, and anything not needed to run Market #1 honestly (see Section 02).

North-star proposition — LOW-TRAINING, NOT ZERO-TRAINING

A seller should complete routine work without formal instruction, recover safely from mistakes, understand what money is available and why, and receive contextual help when the market or task becomes complex.

- **Global by default:** shared workflows and data contracts across markets.
- **Local by configuration:** language, currency, tax, fee, identity, logistics, and payout behavior are versioned by market.
- **Evidence-aware:** behavioral principles guide hypotheses; observed seller performance decides whether a design ships.
- **Safe automation:** financial, identity, and irreversible actions remain reviewable and recoverable.
- **Ruthless v1 scope (new):** every capability is either a named v1 bet, an explicit cut with a revisit condition, or documented debt with an owner and extraction plan.

Section 02 — V1 bets and explicit cuts (NEW)

A blueprint that keeps everything prioritizes nothing. This section is the contract for v1: three bets funded fully, and a cut list that is as binding as the bet list.

The three v1 bets

#	Bet	Why this one	v1 success signal
1	Financial truth: itemized ledger, estimated seller proceeds with rule version, and the estimated/pending/final state model.	Opaque money is the highest-trust failure and the hardest to retrofit.	Sellers correctly explain their payout amount and its uncertainty; estimate-to-final variance is measured with reason codes.
2	Batch fulfillment with safe recovery: multi-select, print-first labels, explicit sync states, reprint history.	Highest-frequency task; largest measurable time savings; failures (duplicate shipments) are costly and visible.	Orders processed per active minute improves with no increase in duplicate shipments or missed cutoffs.
3	Rapid listing with progressive disclosure: essential fields first, conditional market requirements, save-as-draft, honest completeness state.	First-session make-or-break; feeds everything downstream.	Valid publish rate and time-to-first-publish improve; no hidden-requirement surprise at submission.

Explicit v1 cuts

Each cut has a revisit condition. Cutting is not rejecting; it is sequencing.

Cut from v1	Reason	Revisit when
Contribution margin view	Depends on seller-maintained COGS data most solo sellers will not enter; risks being empty or misleading. Estimated proceeds already delivers the trust win.	≥30% of active sellers voluntarily enter cost data, or growing-merchant segment demands it in research.
Predictive prioritization / AI-assisted workflows	Optimization on top of a core that does not exist yet.	Phase 4, after core metrics are stable.
Multi-market configuration console (self-serve UI)	Market #1 needs versioned rules, not a rule-authoring product. Config can live in reviewed files with effective dates.	Second market onboarding begins.
Full adapter marketplace / pluggable everything	One carrier, one payout rail set, one identity provider per launch market is enough to prove the contracts.	Third market, when reuse pressure is real.
Wholesale, pre-order, cross-border listing options	Long-tail listing complexity; keep the schema extensible, hide the UI.	Segment research shows material demand in a launch market.

Advanced analytics dashboards	Operational counts and exceptions suffice for v1.	Growing-merchant/high-volume segments in Phase 3+.
-------------------------------	---	--

RULE: Anything not on the bet list or the cut list is by default *not in v1*. Additions require removing something of equal size.

Section 03 — Scope, evidence, and limitations

This document is a product blueprint built from desk research. It does not claim to represent every seller or market.

Stage 1 field research (Section 13) is a formal go/no-go gate, not a formality: if priority markets contradict the friction hypotheses below, the backlog is re-cut before build continues.

Why the global scope matters

Seller operations look similar at the task level — list products, fulfill orders, manage stock, resolve issues, receive funds — but the surrounding rules vary substantially. Payout timing, currencies, reserves, fees, refunds, and bank processing differ by country, provider, risk level, and transaction type.[2][3][4][5][6] A global seller center cannot hard-code one country's workflow or one flat fee model.

Evidence hierarchy

Level	Used for	Interpretation
Official standards and regulators	Accessibility, authentication, localization, identity, tax examples	Product constraints or strong implementation guidance
Official platform documentation	Fee structures, payout timing, currencies, reserves	Current examples, not universal marketplace rules
Foundational behavioral studies	Hick-Hyman, Fitts, loss aversion, peak-end	Narrow mechanisms; not automatic proof of a UX feature
Original draft and public seller commentary	Candidate pain points	Hypotheses requiring market-specific validation

Known limitations

- The original draft did not provide a transparent interview sample, usability-test protocol, market breakdown, or reproducible click-count study.
- Public help pages describe policy, not necessarily real seller comprehension or performance.
- Named psychological effects have boundary conditions; features still need task testing, comprehension checks, and telemetry.
- Fee and regulatory content changes. Every market rule shown in-product needs an effective date, owner, source, and rollback path.

RESEARCH STATUS: Treat the recommendations as a prioritized design-and-validation backlog. They become evidence-backed decisions only after research with sellers in each launch market.

Section 04 — What changed from the original draft

Original framing	Revised framing	Why
Zero-training Seller Center	Low-training, self-explanatory, recoverable	No complex operational product is literally training-free in every market.
Philippine MSMEs as the design boundary	Global segments; Philippines as one reference market	Core tasks travel; market rules do not.
Real-time net profit	Estimated seller proceeds (margin view deferred — see Section 02)	Payout is not profit; seller costs must be handled separately, and only when sellers actually supply them.
Guaranteed net cash	Estimated / pending / final settlement states	Refunds, reserves, taxes, promotions, FX, and adjustments can change the result.[2][3][4][6]
Automatic daily e-wallet payout	Opt-in schedule and eligible destinations by market	Availability, timing, thresholds, holidays, risk, and regulation vary.[1][2][3]
Peak-End Rule justifies payout automation	Liquidity, predictability, control, reduced manual work	Peak-end research concerns remembered experience, not payout scheduling.[15]
Task cards leverage Zeigarnik + endowed progress	Externalized task state, priority, risk, honest progress	Earned progress is not endowed progress; test without theory theater.
Fitts's Law = automation	Fitts applies to target acquisition only	Large, well-spaced actions are defensible; automatic logistics needs separate safety evidence.[8][13]
Flat 3% + 2% fee example	Versioned itemized fee rules with confidence state	Official documentation shows multiple fee types and regional differences.[4][5]
Everything ships in the blueprint (implicit)	Three named bets; binding cut list; debt register	A plan that funds everything equally funds nothing adequately.

Section 05 — Global seller workflow model

Design around recurring jobs, then configure market-specific policy and integrations.

Seller job	Common friction hypothesis	Target experience	Primary outcome
Create and publish	Long forms, unclear required attributes, repeated data entry	Essential fields first; conditional requirements; save draft; bulk options	Publish accurately with fewer late-stage errors
Fulfill orders	Status hunting, per-order repetition, label/schedule coupling	Batch queue; print-first where permitted; clear carrier cutoff and exceptions	Process more orders with fewer misses
Manage inventory	Deep navigation, variant tables, accidental edits	Inline editing, search/scan, undo, audit trail, low-stock thresholds	Keep stock accurate and recover from mistakes
Serve buyers and resolve issues	Messages detached from orders and dispute deadlines	Order-linked conversation, structured actions, urgency and policy context	Reduce context switching without a noisy inbox
Understand and receive funds	Opaque deductions, estimate/final mismatch, payout uncertainty	Itemized ledger, status model, rule dates, configurable eligible payout schedule	Improve financial comprehension and trust

Seller segments to design for

Segment	Typical operating condition	Design implication
Solo / micro seller	Mobile-first, intermittent operations, low volume	Fast defaults, plain language, contextual guidance, resilient sync
Growing merchant	Multiple channels, moderate volume, small team	Batch actions, roles, audit history, exports, exception queues
High-volume operator	Dedicated staff, scanners/printers, SLA pressure	Keyboard/scanner speed, bulk tools, integrations, operational analytics
Assisted seller	Agent, service center, or shared device	Delegated permissions, privacy, session controls, proof of action

A seller may move between segments by season, category, device, or staffing. Do not equate "small seller" with "simple business." **V1 optimizes for the solo/micro and growing-merchant segments in the first launch market; high-volume and assisted flows are validated but not fully built until Phase 3.**

Section 06 — Product architecture: minimum viable core, deferred by design

Revised: the original architecture list was correct in direction but too heavy as a Phase 1 gate. It is now split into what Market #1 genuinely requires versus what is deliberately deferred with documented debt.

SELLER EXPERIENCE LAYER — Mobile • Web • Assisted operations **GLOBAL CORE** — Task model • Listing • Orders • Inventory • Messaging • Ledger • Permissions • Audit **MARKET CONFIGURATION** — Language • Currency • Fees • Tax • Identity • Payout rules • Service levels (versioned rules with effective dates) **LOCAL ADAPTERS** — Carriers • Banks/wallets • Tax invoices • ID/KYC • Notifications • Support

Tier 1 — Required before Market #1 (non-negotiable)

Capability	Requirement
Versioned fee/tax/policy rules	Every rule shown to a seller has a market, scope, source, owner, effective date, and version. May live in reviewed configuration files — a rule-authoring UI is not required.
Explainable calculation output	Every displayed amount maps to line items and a formula version — not only a single payout number.
Event and audit model	Record who changed stock, payout settings, fulfillment status, and financial inputs; expose human-readable history.
Idempotent actions and sync states	Prevent duplicate submissions; show pending/synced/failed explicitly. Full offline conflict-resolution UI may be simplified to "block and explain" in v1.
Baseline permissions	Owner vs. staff separation protecting payouts, refunds, exports, and account changes. Fine-grained role builder is deferred.

Tier 2 — Deferred with a documented-debt register

Market #1 **may** ship with hard-coded shortcuts in these areas, provided each entry in the debt register has an owner, the assumption written down, and an extraction plan triggered by the second market.

Deferred capability	Acceptable v1 shortcut	Extraction trigger
Generic adapter contracts (carrier, identity, payout, invoice)	Direct integration with one provider per category, behind an internal interface	Second-market kickoff
Multi-market configuration console	Reviewed config files with schema validation	Second-market kickoff
Full offline queue with conflict merge	Read-only offline + queued safe actions + explicit block on conflicts	Field evidence of connectivity-driven task failure
Progressive fine-grained permissions	Two roles (owner, staff) with sensitive-action gating	Growing-merchant/high-volume onboarding

DEBT RULE: Undocumented hard-coding is a defect. Documented, owned, dated hard-coding is a legitimate sequencing decision.

Section 07 — Recommendation set A: action and listing

A1. Action-first overview

DESIGN RULE: Show the smallest set of tasks that require attention now. Each card states the action, quantity, deadline or risk, and next step. Progress must be factual; urgency must not be manipulative.

- Separate required work, risk/exception work, and optional growth opportunities.
- Allow sellers to dismiss or mute non-critical recommendations and promotional content.
- Use counts plus meaning: "5 orders to pack before 16:00" beats "5 pending."
- Provide an "all tasks" view so prioritization does not hide work.
- Expose accessible status updates when counts or sync state change.[8]

A2. Rapid listing with progressive disclosure

Stage	Show by default	Reveal when relevant
Start	Photo/media, product name, price, available quantity	Category suggestions and draft status
Market requirements	Only requirements known from seller/category/market	Tax, safety, brand, customs, or category attributes
Selling options	Shipping profile and lead time	Variants and bundles. <i>(Pre-order, cross-border, wholesale: schema-ready but UI cut from v1 — Section 02.)</i>
Quality and review	Completeness, errors, payout estimate	Search guidance, richer media, optional optimization

- Hick-Hyman research concerns choice reaction time as alternatives increase; it supports simplifying meaningful choices, not blindly hiding fields.[12]
- Never defer a mandatory field until final submission without warning. Show what remains and why it is required.
- Support save-as-draft, duplicate listing, barcode scan, and market templates. Bulk upload ships when the growing-merchant segment is onboarded.
- Measure completed valid listings, correction rate, abandonment point, and time-to-first-publish — not a universal "60-second" promise.

Section 08 — Recommendation set B: fulfillment and inventory

B1. Batch fulfillment with safe automation

Interaction	Recommended behavior	Safety condition
Select orders	Multi-select by status, carrier, cutoff, location, or SKU	Explain excluded orders; preserve selection after errors
Print labels	One primary batch action; print preview and reprint history	Avoid duplicate shipment creation; mark reprints visibly
Arrange pickup / handoff	Auto-suggest or auto-run only when market policy permits	Seller can review schedule, capacity, exceptions, failure state
Update status	Event-driven sync with explicit pending/synced/failed states	Never imply success before carrier confirmation

Fitts's Law supports reducing movement difficulty through larger and closer targets; WCAG 2.2 also establishes minimum target size/spacing guidance.[8][13] These principles justify prominent batch actions, **not** automatic logistics decisions.

B2. Inline inventory controls with recovery

- Direct numeric entry for large changes; stepper controls for small changes.
- Show variant, location, reserved quantity, and sellable quantity distinctly.
- Immediate undo; audit trail with actor, old value, new value, source, and time.
- Require review for unusually large changes, destructive actions, or cross-location overwrites.
- Handle concurrent edits explicitly; do not silently accept the last write. (*V1 may block-and-explain on conflict rather than merge — see Section 06 debt register.*)

ACCESSIBILITY AND SAFETY: Financial and data-changing actions should be reversible, checked, or confirmed, aligned with WCAG 2.2 guidance for preventing serious input errors.[8]

Section 09 — Recommendation set C: financial transparency

Replace false certainty with a clear estimate, an itemized ledger, and a final settlement state.

RETIRE THIS COPY: "Gross price minus fees = guaranteed net cash." It is inaccurate unless the platform contractually absorbs every later tax, refund, reserve, currency, and adjustment difference.

V1 calculation (funded bet)

Measure	Suggested formula	Label and use
Estimated seller proceeds	Item revenue + collected shipping + platform subsidy – platform fees – seller-funded discounts – tax/withholding – known adjustments	Show during pricing and before fulfillment; label as <i>estimated</i> with rule version
Final settlement	Posted credits – posted deductions for the settlement period	Reconciled ledger tied to orders, refunds, reserves, FX, tax, and payout

Deferred calculation (explicit cut — see Section 02)

Measure	Status	Rationale
Estimated contribution margin (proceeds – COGS – packaging – advertising – fulfillment – other seller-entered costs)	Deferred to Phase 4. Ship only if sellers demonstrably maintain cost data. Never infer missing costs as zero.	Depends entirely on voluntary data entry; a mostly-empty or silently-wrong margin number damages the same trust the ledger is meant to build.

Required financial states

Estimated	Pending	Final
Uses current known rules; may change. Show assumptions, missing inputs, rule date, and an uncertainty notice.	In progress. Show holds, reserves, refund windows, payout eligibility, and expected next update.	Posted to the ledger. Show traceable line items, settlement ID, payout status, downloadable records.

- Do not hard-code one fee percentage. Official marketplace pages show multiple fee types, category differences, regulatory fees, reserves, and refund effects.[2][4][5][6]
- Use ISO currency codes where a symbol is ambiguous; format currency, dates, and numbers using locale data such as Unicode CLDR.[11]
- Track estimate-to-final variance and top reason codes. A precise-looking estimate that is frequently wrong damages trust.
- Financial transparency may reduce unpleasant surprise and reference-point mismatch, but loss aversion is not proof this exact interface works.[14]

Section 10 — Recommendation set D: context and payouts

D1. Order-linked conversational hub

- Unify buyer messages, order events, support cases, and dispute alerts around the order — not one undifferentiated stream.
- Structured actions for invoice, refund, discount, address issue, proof upload, tracking update.
- Separate service-level deadlines and legal/policy notices from ordinary chat.
- Preserve search, filters, assignment, internal notes, translation status, audit history.
- Accessible status announcements without excessively interruptive screen-reader output.[8]

D2. Configurable payout experience

Do	Do not
Show eligible payout destinations, currency, timing, thresholds, fees, and settlement prerequisites by market.	Promise instant or daily payout in markets or risk tiers where settlement and bank processing do not support it.[2][3]
Allow opt-in daily, weekly, monthly, threshold-based, or manual schedules where permitted.	Assume e-wallets are universally available, interchangeable with bank accounts, or free.
Require strong re-authentication, notification, and a cooling-off or hold mechanism for destination changes.	Remove identity checks merely to reduce friction; account-change security is not routine payout friction. [9][10]
Show scheduled, processing, sent, deposited, failed, held, and reversed states with trace IDs where available.	Use one success state that implies funds are already visible in the seller's account.

BEHAVIORAL CORRECTION: Do not justify automated payouts with the Peak-End Rule. The defensible product goals are liquidity, predictability, seller control, and fewer manual steps; the original peak-end experiments studied remembered experience.[15]

Section 11 — Global readiness requirements

Localization is product behavior, not translation added at the end.

Dimension	Global-core requirement	Market configuration examples
Language and direction	Externalize all copy; support text expansion; accessible labels	Language, plural rules, LTR/RTL, transliteration, support language
Currency and numbers	Decimal-safe money model; explicit currency codes	Symbol placement, grouping, decimals, accounting negatives, FX disclosure[11]
Time and service levels	Store UTC plus market zone; unambiguous timestamps	Weekends, holidays, carrier cutoffs, dispute and refund windows
Fees, tax and invoices	Versioned line-item ledger and rule engine	VAT/GST/sales tax, withholding, regulatory fee, invoice schema, exemptions
Identity and permissions	Role-based access; recovery and audit	KYC/CDD evidence, beneficial owner, local ID, age and business-status rules[10]
Payouts	Destination abstraction and settlement state model	Bank, card, wallet or mobile-money rails; currency, threshold, schedule, hold
Logistics	Carrier adapter contract; label and manifest model	Pickup/drop-off, address format, dimensions, COD, reverse logistics, customs
Connectivity	Idempotent actions; sync queue; explicit conflict handling	Low-bandwidth mode, shared devices, scanner/printer availability
Accessibility	WCAG 2.2 AA target; keyboard, screen reader, zoom, color	Local assistive-tech behavior and translated accessibility QA[8]

PHILIPPINES IN THE GLOBAL MODEL: Philippine language, tax/withholding, GCash/Maya eligibility, local banks, courier rules, COD, address conventions, and connectivity conditions become one launch-market configuration — not assumptions embedded in the shared product core.[7]

(Per Section 06, generic adapter contracts and the configuration console are Tier 2: Market #1 may use direct, documented integrations.)

Section 12 — Behavioral design: keep, reframe, or retire

Named effects should clarify a mechanism, never substitute for user evidence.

Principle / label	Decision	Correct use in this product
Progressive disclosure	KEEP	Stage complexity while preserving visibility of mandatory requirements and completion status.
Hick-Hyman Law	NARROW	Use when reducing or structuring meaningful alternatives; do not equate every long form with Hick's Law.[12]
Fitts's Law	KEEP	Target size, spacing, and distance in repeated actions; never to justify automation.[13]
Loss aversion	REFRAME	A caution about unexpected deductions and reference points; validate financial comprehension directly.[14]
Peak-End Rule	RETIRE FOR PAYOUTS	Does not establish that daily automated payout is the right schedule.[15]
Zeigarnik Effect	REFRAME	Honest task visibility and externalized memory; do not rely on unfinished-task tension as the success mechanism.
Endowed progress	REMOVE LABEL	A truthful "3 of 5 packed" indicator is earned progress, not endowed progress.
System 1 automation	REMOVE LABEL	Describe reduced choices, learned routines, defaults, and error recovery instead of a broad dual-process claim.
Mental context	KEEP WITHOUT THEATER	Link messages, order state, deadlines, and actions so sellers do not reconstruct context.

Evidence gate for every behavioral claim

1. State the mechanism precisely: what user behavior should change, and why?
2. Define a comparison: current experience, alternative design, or baseline task.
3. Measure task outcome, errors, comprehension, confidence, and downstream operational impact.
4. Check for harm: pressure, accidental action, hidden requirements, exclusion, or false certainty.
5. Keep the feature only if observed outcomes justify it; the theory name is optional.

Section 13 — Research and validation plan

Validate the global core across seller segments, then validate each market configuration before launch.

Stage 1 — Discovery and workflow evidence (GO/NO-GO GATE)

Revised: Stage 1 is now a formal gate. Its explicit purpose is to confirm or kill the desk-research hypotheses in Sections 05–10 before core build continues past the foundation phase.

- Select 4–6 priority markets differing in language, payment rails, logistics maturity, regulatory context, and connectivity.
- Recruit solo/mobile-first, growing-team, and high-volume sellers; include assisted/shared-device operations where relevant.
- Use contextual interviews, screen walkthroughs, artifact review, and task diaries. Record platform, category, volume, staff, devices, fulfillment model.
- Map actual exceptions: failed pickup, refund, partial fulfillment, inventory conflict, identity review, payout hold, account change.
- **Gate output:** each friction hypothesis is marked confirmed, revised, or killed, with evidence. A killed hypothesis removes its feature from the bet list; the v1 cut list (Section 02) is re-balanced accordingly.

Stage 2 — Iterative prototype testing

Task	Critical observation	Pass signal
Publish a product	Required-field comprehension and recovery	Valid listing without hidden-requirement surprise
Print and hand off a batch	Selection, exceptions, reprint and schedule understanding	No duplicate shipment; exceptions resolved
Correct inventory	Precision, accidental change and recovery	Correct quantity; undo/audit understood
Explain expected payout	Difference among revenue, proceeds, and final settlement	Seller accurately explains amount and uncertainty
Change payout destination	Security, delay and confirmation comprehension	Authorized change; seller understands hold and notification
Recover from offline failure	Duplicate action and conflict handling	No data loss or silent overwrite

Stage 3 — Market pilot

- Run an instrumented pilot long enough to include at least one payout cycle and common refund/return windows.
- Review by market, segment, device, language, accessibility need, and order volume; do not hide local failure inside a global average.
- Maintain a research repository with evidence, decision, owner, confidence, and date. Re-test when market policy or integrations change materially.

Section 14 — Measurement framework and release gates

Measure operational success and financial comprehension, not clicks or engagement alone.

Outcome	Primary measures	Guardrail
Listing quality	Valid publish rate; time to valid listing; correction and abandonment point	Policy rejection, duplicate listing, support contact
Fulfillment speed	Orders processed per active minute; batch completion; exception resolution	Missed cutoff, wrong label, duplicate shipment
Inventory accuracy	Successful change; reconciliation variance; recovery time	Accidental edit, oversell, unresolved sync conflict
Financial clarity	Payout-comprehension task; estimate-to-final variance; reason-code coverage	Misleading certainty, unexplained adjustment, dispute volume
Payout reliability	On-time rate by state; failure recovery; traceability	Unauthorized destination change, hidden hold, notification failure
Accessibility	Task completion with keyboard/screen reader/zoom; issue severity	Any blocker in a core seller task
Global quality	Locale defect rate; translation comprehension; market-config accuracy	Wrong currency/tax/time, RTL breakage, stale policy version

Release gates

Gate	Required evidence
Core usability	Priority tasks meet agreed thresholds in each launch segment; no critical recovery gap.
Financial truth	All displayed amounts map to ledger line items and a versioned rule; estimate status and variance reasons visible.
Market readiness	Localized rules, copy, currency, time, identity, logistics, and payout adapters pass market-owner review.
Security and permissions	Sensitive changes use appropriate authentication, notification, audit, and least-privilege access.[9][10]
Accessibility	Core flows meet WCAG 2.2 AA acceptance criteria and assistive-technology testing.[8]
Operational recovery	Failed sync, carrier errors, payout failures, and outages have explicit seller-facing states and support paths.
Debt register (new)	Every Tier 2 shortcut is documented with owner, assumption, and extraction trigger; no undocumented hard-coding in shipped code.

Section 15 — Implementation roadmap

Revised: Phase 1 is re-scoped around the three v1 bets and Tier 1 architecture only. The original 8–12 week Phase 1 assumed Tier 2 capabilities that are now deferred; with the narrower scope the window is honest rather than optimistic.

Phase	Focus	Deliverables	Exit condition
0. Evidence and policy foundation — 4–6 weeks	Stage 1 discovery (go/no-go); event model; fee/policy schema; security and accessibility requirements	Confirmed/revised/killed hypothesis register; market rules documented with owners and confidence	Stage 1 gate passed; bet list re-confirmed
1. Trustworthy operating core — 10–14 weeks	The three bets only: rapid listing; batch fulfillment; ledger + estimated proceeds; plus Tier 1 architecture (versioned rules, explainable amounts, audit, idempotency, baseline permissions)	Core tasks pass usability, financial-truth, accessibility, and recovery gates	Bets meet their Section 02 success signals in prototype testing
2. First market launch — 6–10 weeks	One carrier, one payout rail set, one identity provider, tax/invoice, language, support — direct integrations with debt register entries	End-to-end pilot covering settlement, returns, failures, support operations	Pilot spans ≥ 1 payout cycle and refund window; debt register complete
3. Second/third market scale — per market	Extract Tier 2: adapter contracts, configuration console, localization QA, fine-grained permissions	Reusable adapter toolkit; market dashboards	New market ships without forking the global core
4. Optimization — ongoing	Contribution margin (if cost-data condition met), predictive prioritization, bulk tools, assisted workflows, carefully evaluated AI support	Measured improvements	No increase in error, pressure, exclusion, or estimate variance

Product principles to protect during delivery

1. **Truth before delight.** Amounts, states, deadlines, and automation must accurately describe what the system knows.
2. **Fast routine, visible exception.** Optimize common work without hiding uncommon cases that create loss or support burden.
3. **Local rules, coherent product.** Configure differences; do not fork country products unless law or operations require it.
4. **Recovery is part of usability.** Undo, audit, retry, reconciliation, and support paths are first-class interactions.
5. **Evidence over theory theater.** Behavioral science forms testable hypotheses; observed seller outcomes decide what ships.
6. **Scope is a promise (new).** The bet list and cut list are maintained together; additions require equal-size removals, and shortcuts live in the debt register or not at all.

Section 16 — References and source notes

Official and foundational sources used to correct the original draft. Accessed 6 August 2026.

1. World Bank. Global Findex Database 2025. <https://www.worldbank.org/en/publication/globalfindex/report>
2. Shopify Help Center. Understanding Shopify Payments payout timing. <https://help.shopify.com/en/manual/payments/shopify-payments/payouts/payout-timing>.
3. Shopify Help Center. Supported payout currencies and multi-currency payouts. <https://help.shopify.com/en/manual/payments/shopify-payments/store-currency/supported-payout-currencies>
4. TikTok Shop Seller Center Philippines. Introduction to Seller Fees. https://seller-ph.tiktok.com/university/essay?knowledge_id=2675772847064834&lang=en
5. Etsy Help. Fees and Taxes for Selling on Etsy. <https://help.etsy.com/hc/en-us/articles/115014483627-What-are-the-Fees-and-Taxes-for-Selling-on-Etsy>
6. Amazon. How Amazon seller payments work. <https://sell.amazon.com/blog/amazon-seller-payments>
7. Philippine Bureau of Internal Revenue. Revenue Regulations No. 16-2023. <https://bir-cdn.bir.gov.ph/BIR/pdf/RR%2016-2023.pdf>
8. W3C Web Accessibility Initiative. Web Content Accessibility Guidelines (WCAG) 2.2. <https://www.w3.org/TR/WCAG22/>
9. US National Institute of Standards and Technology. Digital Identity Guidelines: Authentication and Authenticator Management (SP 800-63B-4). <https://pages.nist.gov/800-63-4/sp800-63b.html>
10. Financial Action Task Force. Guidance on Digital Identity. <https://www.fatf-gafi.org/en/publications/FinancialInclusionandnpoissues/Digital-identity-guidance.html>
11. Unicode Consortium. CLDR number and currency patterns. <https://cldr.unicode.org/translation/number-currency-formats/number-currency-patterns>
12. W. E. Hick. On the Rate of Gain of Information (1952). <https://doi.org/10.1080/17470215208416600>
13. P. M. Fitts. The Information Capacity of the Human Motor System in Controlling the Amplitude of Movement (1954). <https://doi.org/10.1037/h0055392>
14. A. Tversky and D. Kahneman. Loss Aversion in Riskless Choice: A Reference-Dependent Model (1991). <https://academic.oup.com/qje/article-abstract/106/4/1039/1873382>
15. D. Kahneman, B. Fredrickson, C. Schreiber, and D. Redelmeier. When More Pain Is Preferred to Less: Adding a Better End (1993). <https://journals.sagepub.com/doi/abs/10.1111/j.1467-9280.1993.tb00589.x>

SOURCE NOTE: Platform help pages are cited as current examples of regional and policy variability, not as endorsements or universal benchmarks. Regulatory and platform rules must be re-verified before each market launch.

REVISION NOTE (v2): Changes from v1 review — (a) new Section 02 with three funded bets and a binding cut list; (b) Section 06 architecture split into Tier 1 (required for Market #1) and Tier 2 (deferred with a documented-debt register); (c) contribution margin deferred to Phase 4 with an explicit data-availability condition; (d) Stage 1 discovery converted into a go/no-go gate that can remove features from the bet list; (e) Phase 1 window adjusted to 10–14 weeks against the narrower scope; (f) new debt-register release gate and "scope is a promise" delivery principle.